

Regression Analysis

S. R. M. Prasanna

Dept of Electrical Engineering
Indian Institute of Technology Dharwad

prasanna@iitdh.ac.in

November 23, 2020

Acknowledgements

- Artificial Neural Networks by Prof. B. Yegnanarayana, PHI, 1999
- Machine learning video lectures by Prof. Andrew Ng, Stanford Uty
- Jason Brownlee "Basics of Linear Algebra for Machine Learning", Online book, 2018.

Terminologies in Regression

- **Regression:** Finding mapping function between given input variable(s) and **output variable which is continuous**.
- Mathematically, given $(x, y) = \{(x^{(i)}, y^{(i)})_{i=1,2,\dots,M}\}$, find mapping function $h_w()$, such that $y = h_w(x)$.
- Using $h_w()$ for unseen data, $\hat{y}_k = h_w(x_k)$, $k \neq i$ and $\hat{y}_k \approx y_k$.
- **Univariate vs Multivariate:** x is 1D vs 2D or more.
- Univariate implies regression with one variable. $y = w_0 + w_1x$
- Multivariate $\implies y = w_0 + w_1x_1^{(i)} + w_2x_2^{(i)}$
- **Linear vs non-linear:** coefficients are linear vs nonlinear.
- **Univariate Linear Regression:** $y = w_0 + w_1x$
- **Univariate Non-Linear Regression:** $y = w_0 + (w_1)^2x$ or $y = \exp(w_0 + w_1x)$

Univariate Linear Regression

- $y = h_w(x) = w_0 + w_1x$
- Need to estimate w_0, w_1 which are best fit for the given data (x, y) .
- Estimate of y : $\hat{y}_i = h_w(x_i)$.
- Error: $e_i = \hat{y}_i - y_i$.
- Squared Error: $e_i^2 = (\hat{y}_i - y_i)^2$.
- Mean Squared Error (MSE): $E = \frac{1}{M} \sum_{i=1}^M e_i^2 = (\hat{y}_i - y_i)^2$
- Finally, (w_0, w_1) that result in least MSE are chosen as parameters.

Cost Function (C)

- $\min_{w_0, w_1} C(w_0, w_1) = \min_{w_0, w_1} \frac{1}{2M} \sum_{i=1}^M (w_0 + w_1 x_i - y_i)^2.$
- How to minimize $C(w_0, w_1)$?
- Wrt to single variable, say w_1 , cost function i.e., error plot is a **parabola**
- Wrt to two variables, cost function is a **contour** in 2D
- Wrt to three variables, cost function is a **surface** in 3D
- All of them will be convex in nature
- Min point in convex surface will give optimal values for (w_0, w_1)

Gradient Descent Method for Minimization

- More generic method for minimization
- Given function $G(\theta_0, \theta_1)$
- Start with some (θ_0, θ_1)
- Iteratively change (θ_0, θ_1) to minimize $G(\theta_0, \theta_1)$ until minima is reached.
- Values of (θ_0, θ_1) at minima are chosen as optimal parameters.

Gradient Descent Algorithm

- $\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} G(\theta_0, \theta_1), j = 0, 1$
- α is a learning rate parameter.
- Partial derivative essentially represents slope
- $\theta_0 := \theta_0 - \alpha \frac{\partial}{\partial \theta_0} G(\theta_0, \theta_1)$
- $\theta_1 := \theta_1 - \alpha \frac{\partial}{\partial \theta_1} G(\theta_0, \theta_1)$
- Repeat the above two steps until convergence
- Smaller α , smaller step size
- Bigger α , bigger step size
- Bigger step size may lead to non-convergence to minima.
- Also large α in some case may diverge.

Gradient Descent for Linear Regression

- **Univariate LR:** $\min_{w_0, w_1} C(w_0, w_1) = \min_{w_0, w_1} \frac{1}{2M} \sum_{i=1}^M (w_0 + w_1 x_i - y_i)^2.$
- **Gradient Descent:** $G(\theta_0, \theta_1)$
- $\theta_0 := \theta_0 - \alpha \frac{\partial}{\partial \theta_0} G(\theta_0, \theta_1)$
- $\theta_1 := \theta_1 - \alpha \frac{\partial}{\partial \theta_1} G(\theta_0, \theta_1)$
- Repeat the above two steps until convergence
- $\theta_0 = w_0$, $\theta_1 = w_1$, and $G(\theta_0, \theta_1) = C(w_0, w_1)$

Gradient Descent for Linear Regression

- $\frac{\partial}{\partial w_0} C(w_0, w_1) = \frac{\partial}{\partial w_0} \frac{1}{2M} \sum_{i=1}^M (w_0 + w_1 x_i - y_i)^2$
- $\frac{\partial}{\partial w_0} C(w_0, w_1) = \frac{1}{M} \sum_{i=1}^M (w_0 + w_1 x_i - y_i)$
- $\frac{\partial}{\partial w_1} C(w_0, w_1) = \frac{\partial}{\partial w_1} \frac{1}{2M} \sum_{i=1}^M (w_0 + w_1 x_i - y_i)^2$
- $\frac{\partial}{\partial w_1} C(w_0, w_1) = \frac{1}{M} \sum_{i=1}^M (w_0 + w_1 x_i - y_i) x_i$
- w_0, w_1 by Gradient Descent:
- $w_0 := w_0 - \alpha \frac{1}{M} \sum_{i=1}^M (w_0 + w_1 x_i - y_i)$
- $w_1 := w_1 - \alpha \frac{1}{M} \sum_{i=1}^M (w_0 + w_1 x_i - y_i) x_i$
- Repeat the above two steps until convergence

Salary Prediction by Univariate LR

- Let x is experience in months and y is salary in rupees.
- Given $M = 100$ values of (X, Y) .
- Learn mapping between salary (y) and experience in month (x):
 $y = h_w(x)$.
- Prediction of salary for a new x_k using using: $\hat{y}(k) = h_w(x_k)$.
- Expect $\hat{y}(k) \approx y_k$
- Can I better fit the model? Yes, using higher order polynomial in x .
- $y = w_0 + w_1x + w_2x^2 + w_3x^3$

Motivation for Multivariate Linear Regression

- Consider salary prediction problem discussed earlier.
- Let x represent total experience.
- Of course, can better fit the curve using higher order polynomial of x
- However, experience has many components like teaching (x_1), research (x_2), admin (x_3) and so on.
- Thus, in univariate LR $x = \{x_1 + x_2 + x_3\}$.
- Each component represent different aspect. Can I do better prediction using different components explicitly ?
- Hence the motivation for multivariate linear regression.

Multivariate Linear Regression

- **Univariate LR:** $y = h_w(x) = w_0 + w_1x$
- **Multivariate LR:** $y = h_w(x) = w_0 + w_1x_1 + w_2x_2 + \dots + w_Nx_N$
- **Multivariate LR:** $y = h_w(x) = w_0x_0 + w_1x_1 + w_2x_2 + \dots + w_Nx_N$,
where $x_0 = 1$
- $h_w(x) = w^T x$, inner product of transpose of parameter vector with input vector
- $w = \{w_0, w_1, \dots, w_N\}^T$ and $x = \{x_0, x_1, \dots, x_N\}^T$, w, x , each is col vector of $(N + 1)$ dimension

Multivariate Linear Regression using Gradient Descent

- (X, Y) , where, $X = \{x_j^{(i)} | i = 1, 2, \dots, M, j = 0, 1, 2, \dots, N\}$, X is $M \times (N + 1)$ matrix and $Y = \{y^{(1)}, y^{(2)}, \dots, y^{(M)}\}$ is M dim col vector.
- Estimate $w = \{w_0, w_1, w_2, \dots, w_N\}$ which are best fit for (X, Y) .
- $$\min_{w_0, w_1, \dots, w_n} C(w_0, w_1, \dots, w_N) = \min_{w_0, w_1, \dots, w_n} \frac{1}{2M} \sum_{i=1}^M (\hat{y}^{(i)} - y^{(i)})^2.$$
- $$\min_w C(w) = \min_w \frac{1}{2M} \sum_{i=1}^M (w^T x^{(i)} - y^{(i)})^2.$$

Gradient Descent for Linear Regression

- **Gradient Descent:** $G(\theta_0, \theta_1, \dots, \theta_n) = G(\theta)$
- $\theta_0 := \theta_0 - \alpha \frac{\partial}{\partial \theta_0} G(\theta)$
- $\theta_1 := \theta_1 - \alpha \frac{\partial}{\partial \theta_1} G(\theta)$
- $\theta_N := \theta_N - \alpha \frac{\partial}{\partial \theta_N} G(\theta)$
- Repeat the above N steps until convergence
- $\theta_0 = w_0$, $\theta_1 = w_1$, and $\theta_N = w_N$ $G(\theta) = C(w)$

Gradient Descent for Linear Regression

- $\frac{\partial}{\partial w_0} C(w) = \frac{\partial}{\partial w_0} \frac{1}{2M} \sum_{i=1}^M (w^T x^{(i)} - y^{(i)})^2$
- $\frac{\partial}{\partial w_0} C(w) = \frac{1}{M} \sum_{i=1}^M (w^T x^{(i)} - y^{(i)})$
- $\frac{\partial}{\partial w_1} C(w) = \frac{\partial}{\partial w_1} \frac{1}{2M} \sum_{i=1}^M (w^T x^{(i)} - y^{(i)})^2$
- $\frac{\partial}{\partial w_1} C(w) = \frac{1}{M} \sum_{i=1}^M (w^T x^{(i)} - y^{(i)}) x_1^{(i)}$
- $\frac{\partial}{\partial w_N} C(w) = \frac{1}{M} \sum_{i=1}^M (w^T x^{(i)} - y^{(i)}) x_N^{(i)}$
- **w by Gradient Descent:**
- $w_0 := w_0 - \alpha \frac{1}{M} \sum_{i=1}^M (w_0 + w_1 x^{(i)} - y^{(i)}) x_0^{(i)}$
- $w_1 := w_1 - \alpha \frac{1}{M} \sum_{i=1}^M (w_0 + w_1 x^{(i)} - y^{(i)}) x_1^{(i)}$
- $w_N := w_N - \alpha \frac{1}{M} \sum_{i=1}^M (w^T x^{(i)} - y^{(i)}) x_N^{(i)}$
- Repeat the above $(N + 1)$ steps until convergence

Salary Prediction by Multivariate LR

- Let $x = \{x_1, x_2, x_3\}$ represent different experience in months and y represent salary in rupees.
- Given $M = 100$ values of (X, Y) , X is a matrix.
- Learn mapping between salary (y) and experience in month (x):
 $y = h_w(x)$.
- Prediction of salary for a new x_k using using: $\hat{y}(k) = h_w(x_k)$.
- Expect $\hat{y}(k) \approx y_k$
- Can I better fit the model? Yes, using higher order polynomial in x .
- $y = h_w(x) = w_0 + w_1x_1 + w_2x_2 + w_3x_3 + w_4x_4 + \dots$

Parameter Estimation using Normal Equations Process

- $y = Xw$
- $e = Xw - y$
- $E = (Xw - y)^T(Xw - y)$
- $E = ((Xw)^T - y^T)(Xw - y)$
- $E = (Xw)^T Xw - y^T Xw - (Xw)^T y + y^T y$
- $E = (Xw)^T Xw - 2(Xw)^T y + y^T y$
- $E = w^T X^T Xw - 2w^T X^T y + y^T y$
- $\frac{\partial}{\partial w} E = 0$
- $\frac{\partial}{\partial w} E = 2X^T Xw - 2X^T y = 0$
- $X^T Xw = X^T y$
- $w = (X^T X)^{-1} X^T y$

Linear Regression using Normal Equations

- **Normal equations** to find w , than iterative in Gradient descent.
- Method to solve equations rather than iterative procedure.
- **Solving Normal Equations:** Partially derivative of E or $C(w)$ with each dimension of w equating the resulting derivative eqn to zero.
- Results in $(N + 1)$ normal equations.
- **Feature matrix:** X , where each row is $N + 1$ feature vector. There are M such feature vectors.
- y is M dimension output vector and $y = Xw \implies X^T y = X^T X w$
- $w = (X^T X)^{-1} X^T y$.

LR: Gradient Descent(GD) vs Normal Equations (NE)

- GD requires choosing α where as NE does not need.
- GD is iterative vs NE is one step.
- GD works well even when N , feature dim is very large.
- NE becomes difficult due to difficulty in finding $(X^T X)^{-1}$.
- If feature dim is small to moderate dim then use NE.
- If feature dim is very large dim then use GD.